

# Apache Kafka

History, Pub/Sub messaging, queue systems, Kafka architecture, producers, consumers, topics, ISR

Sairam Elangovan

# History

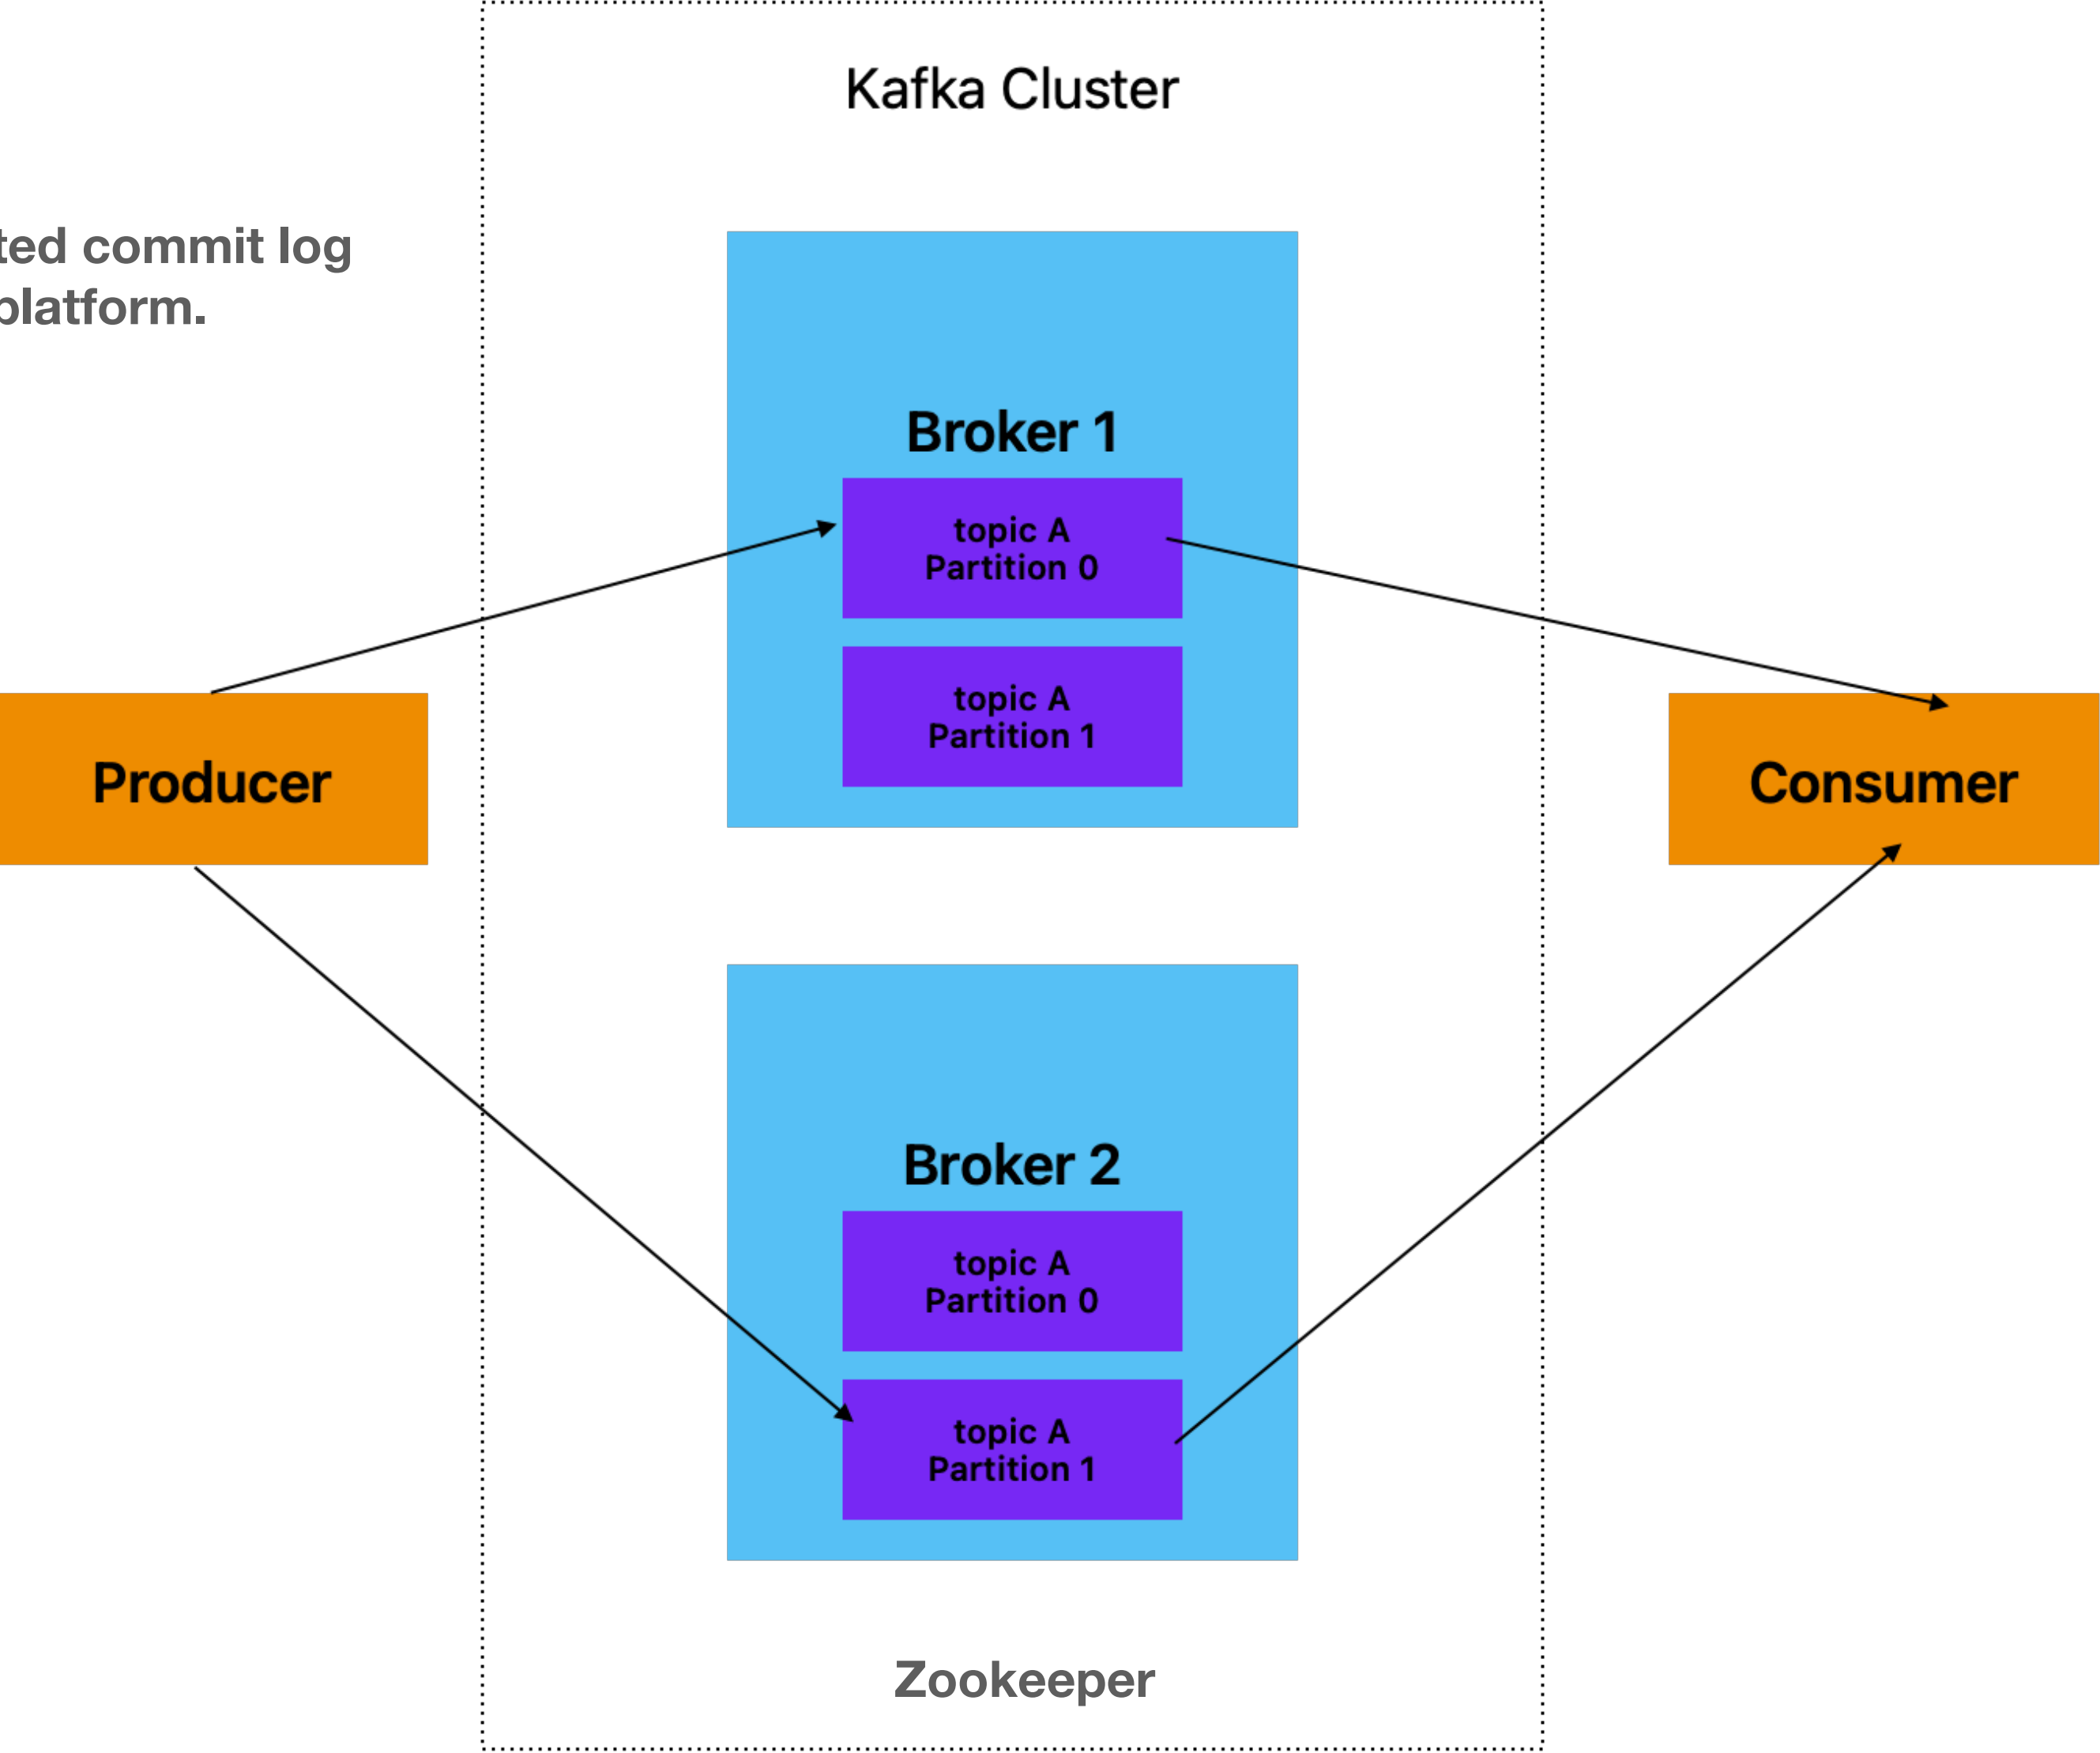
- LinkedIn's system for collecting system and application metrics had performance issues and the data collection was inconsistent across different systems.
- The systems would break continuously due to the ever changing schemas and tracking was based on hourly batching, i.e. it could not be used in real time.

# Primary goals of Kafka

- Decoupled messaging system (push-pull model using producers and consumers)
- Data persistence (multiple consumers consuming the same data for different purposes)
- High Throughput
- Horizontal scaling of system

# Kafka Architecture

Kafka is often described as distributed commit log  
or as a distributed streaming platform.



# Kafka installation

- Tarball local development installation
- Installation via package managers (apt, yum, rpm etc)
- Docker containers

# Kafka installation through docker compose

- `git clone https://github.com/confluentinc/cp-all-in-one`
- `cd cp-all-in-one`
- `cd cp-all-in-one/`
- `docker-compose up -d zookeeper broker control-center`

# Tarball local development installation

- `$ curl -O http://packages.confluent.io/archive/7.3/confluent-7.3.1.tar.gz`
- `$ tar xzf confluent-7.3.1.tar.gz`
- Configuring `.bashrc` file in linux for enabling custom env variable cods (Optional):
  - `$ export CONFLUENT_HOME=${HOME}/confluent-7.3.1 \ && echo "export CONFLUENT_HOME=$CONFLUENT_HOME" >> ~/.bashrc`
  - `$ echo "export PATH=$CONFLUENT_HOME/bin:${PATH}" >> ~/.bashrc`
  - `$ ~/confluent-7.3.1/bin/confluent completion bash | sudo tee /etc/bash_completion.d/confluent \ && echo "source /etc/bash_completion.d/confluent" >> ~/.bashrc \ && source ~/.bashrc`

# Kafka installation via package managers

- `$ sudo apt-get install openjdk-11-jre-headless`
- `$ wget -qO - https://packages.confluent.io/deb/7.1/archive.key | sudo apt-key add -`
- `$ sudo add-apt-repository \ "deb [arch=amd64] https://packages.confluent.io/deb/7.1 stable main" && \ sudo apt-get update`
- `$ sudo apt-get install -y \ confluent-platform \ confluent-security`
- `$ sudo systemctl start confluent-zookeeper` (systemctl enable for running services even if your linux server is rebooted)
- `$ sudo systemctl start confluent-server`



# Zookeeper's role in kafka

- ZooKeeper is a centralized service for maintaining configuration information, naming, providing distributed synchronization, and providing group services.
- All of these kinds of services are used in some form or another by distributed applications.
- ZooKeeper is used for
  - Controller election
  - Cluster membership
  - Topic configuration
  - Access control lists
  - Quotas
  - Maintaining consumer offsets ( post 0.9v consumer offsets are stored under \_\_consumer\_offsets topic)

# Kafka brokers

- A computer instance or container running Kafka process
- Manage partitions
- Handle read and write requests
- Manage replication of partitions

# Topic creation

- Topics are logs that hold messages or events in a logical order.
- ISR- All reads/writes for a specific partition happens through 'Leader' of the partition and 'Follower' get in-sync with 'Leader' for updates.
- `$ bin/kafka-topics.sh --create --topic new-topic --bootstrap-server localhost:9092`
- `$ bin/kafka-topics.sh --describe --topic quickstart-events --bootstrap-server localhost:9092`
- `bin/kafka-topics.sh --list --zookeeper localhost:2181`

# Producers

- A producer partitioner maps each message to a topic partition, and the producer sends a produce request to the leader of that partition
- The partitioners shipped with Kafka guarantee that all messages with the same non-empty key will be sent to the same partition.
- `$ bin/kafka-console-producer.sh --topic quickstart-events --bootstrap-server localhost:9092`

# Consumers

- Consumers read data from Kafka topics.
- Consumers can be standalone or a part of consumer group
- Offset management ( from-beginning, latest and none)
- `$ bin/kafka-console-consumer.sh --topic quickstart-events --from-beginning --bootstrap-server localhost:9092`

# Docker exec commands

- `docker exec broker \ kafka-topics --bootstrap-server broker:9092 \ --create \ --topic new-topic`
- `docker exec --interactive --tty broker kafka-console-producer --bootstrap-server broker:9092 --topic new-topic`
- `docker exec --interactive --tty broker kafka-console-consumer --bootstrap-server broker:9092 --topic new-topic --from-beginning`
- `docker exec broker kafka-topics --bootstrap-server broker:9092 --list`